

FL04 Ship an Automation Workflow v2


WORKFLOW: AUTOMATED COMPANY RESEARCH AND OUTREACH PREPARATION PIPELINE


1. Step Diagram

The workflow I designed was a seven-step company research and drafting pipeline. The objective was to take a new company as the input and systematically move from current company research through to a research-backed output suitable for identifying relevant people, hiring signals, technical direction, opportunities, and potential outreach.

The intended workflow was:

NEW COMPANY INPUT





Each step was deliberately separated rather than asking one model prompt to perform the entire task at once. This created explicit handoffs between research, analysis, synthesis and drafting.

Why I used this structure

The most important requirement for this workflow was **current and accurate information**. Company research depends heavily on live information such as websites, LinkedIn profiles, current job postings, recent announcements, product direction and other public signals.

I initially attempted to build the workflow in **n8n**, including the required nodes and connections. However, finding a free API/model combination that could reliably support the complete workflow proved unsuccessful.

I was able to set up the nodes and workflow structure, but the model that I could get working end-to-end did not have the required web-search capability. This made it unsuitable for this particular use case because a company-research workflow that cannot access current information risks producing outdated or inaccurate results.

I therefore moved to a **clear separation of the seven tasks as individual structured prompts in ChatGPT**, where web search could be used when required. This preserved the intended workflow logic while prioritising the accuracy and recency of the research.

2. Five Real Runs

The workflow was run against five real companies rather than only testing it on a synthetic example.

Run 1 — Ontora | Input: Ontora | **Website:** ontora.com

Result:

FINAL OUTREACH 1

Hi Leon,

Ontora's core loop is interesting because the difficult part isn't getting an LLM to summarize an interview; it's deciding whether the operational knowledge extracted from employee interviews is trustworthy enough to become a context layer other systems can act on.

I'd like to try to break that part.

If you gave me a small set of anonymized or synthetic interview transcripts, I could build an adversarial test set around contradictory employee accounts, omitted facts, incorrect attribution and plausible-but-unsupported conclusions, then report back where the generated operational picture agrees with ground truth and where it doesn't.

I've done a similar kind of validation on my own systems: for an agentic BI system I built a fault-injection suite that caught 24/24 deliberately introduced issues, and on an ML project I found critical data leakage behind an apparently strong result.

I'm not assuming Ontora has this problem. The point would be to test it rather than speculate.

Would it be useful to spend 20 minutes picking a small test case?

FINAL STRATEGY 2

GraphRAG disagreement and grounding evaluation.

FINAL OUTREACH 2

Hi David,

One part of Ontora's technical direction caught my attention: the GraphRAG layer is designed to answer questions across employee interviews while returning citations, patterns and divergent opinions.

I'd be curious how it behaves when two employees describe the same process differently and both accounts are locally plausible.

I could build a small adversarial corpus with deliberately conflicting accounts and test whether the resulting answers distinguish consensus from disagreement, preserve attribution and avoid turning an unsupported synthesis into a fact.

My relevant experience is building a LangGraph system with reflection/retry and then stress-testing it with 24 injected failures, alongside ML validation work where I had to audit apparently strong results rather than assume they were correct.

If you've already got this evaluation covered, that's useful to know. If not, I could make the test concrete rather than send another generic AI pitch.

Worth a short conversation?

FINAL STRATEGY 3

API/agent reliability testing.

FINAL OUTREACH 3

Hi Max,

Ontora now exposes the product through a REST API, webhooks, CLI and MCP, which makes predictable behavior important once the AI output is being consumed by other software.

I'd be interested in taking one representative workflow and building a small reliability test around it — malformed structured outputs, retries, missing fields, downstream validation and other edge cases around the happy path.

I've worked across that integration boundary before: Flask/FastAPI services, authentication, rate limiting, Pytest and ML/API integration, including taking over backend integration on SolarSense when the assigned backend developer stepped back.

I wouldn't assume there's a gap in Ontora's current testing. I'd rather test a small slice and see whether anything useful falls out.

Open to a short call to choose the workflow?

RECOMMENDED OPTION

FINAL STRATEGY 1 — adversarial evaluation of the core interview-to-knowledge pipeline.

WHY THIS OPTION

It has the best combination of:

- direct relevance to Ontora's central product;*
- strong evidence from Ontora's public architecture;*
- strong evidence from Theresa's actual experience;*
- bounded trial scope;*
- low dependence on an unverified technical assumption.*

It also avoids the weakest form of cold outreach: claiming to know what is wrong with the company's system without having seen it.

HUMAN REVIEW REQUIRED

Before sending, Theresa should verify:

- the company need is actually worth pursuing;*
- the contact is the right person;*
- the email address is correct;*
- the proposed intervention is realistic;*
- the message sounds like her;*
- she is comfortable sending it.*

No prospect outreach has been sent.

Run 2 — Autostep | Input: Autostep | **Website:** autostep.ai

Result:

FINAL OUTREACH 1

Hi Aidan,

Autostep's approach of clustering observed workflow data to figure out which work is actually worth automating struck me as an interesting evaluation problem: the difficult cases aren't necessarily obvious failures, but tasks that look repetitive while meaning something different in context.

I'd like to try to break that layer rather than assume it's working. If you have a small anonymized or synthetic slice of workflow data, I could build a focused evaluation around ambiguous task boundaries, noisy observations and misleading repetition signals, then report back where the classifications hold up and where they don't.

I did a similar kind of validation in my own ML work: an apparently excellent model result turned out to be data leakage, and in another system I deliberately injected 24 issues and caught all 24 through testing.

If you've already covered this internally, no problem. If not, would you be open to a short call to see whether a small test would be useful?

FINAL STRATEGY 2

Independent validation of Human ROI / Agent ROI measurement.

FINAL OUTREACH 2

Hi Aidan,

Autostep's recent positioning around Human ROI and Agent ROI seems like an important distinction: finding work that can be automated is one question; showing that the intervention actually improved the work is another.

I'd be interested in testing that measurement approach on a small example. Given a before/after workflow, I could structure an evaluation that separates time/cost savings from output quality and agent/tool cost, so the result isn't simply that an agent completed the task.

A lot of my ML work has involved selecting and validating systems against the metric that actually mattered—for example, choosing a churn model based on recall rather than raw accuracy.

If you already have this evaluation framework nailed down, I'd leave it alone. Otherwise, would a short conversation about one real workflow make sense?

FINAL STRATEGY 3

Adversarial testing of agent/MCP workflows.

FINAL OUTREACH 3

Hi Aidan,

I noticed Autostep is pushing further into MCP usage, context mining and agentic workflows. That introduces failure modes beyond ordinary model errors: wrong tool selection, incomplete context, unnecessary calls, or poor recovery from malformed results.

I'd like to run a small adversarial test rather than speculate about where those failures might be. Give me a controlled or synthetic workflow and I can deliberately introduce a few of those conditions, measure what the system catches, and produce a short failure taxonomy.

I built a LangGraph system with reflection and retry and stress-tested it by injecting 24 issues, catching all 24. I wouldn't assume that maps directly to Autostep, but the testing pattern is transferable.

If your existing testing already covers this, no issue. If there is a gap, would you be open to a short call to pick one workflow to test?

RECOMMENDED OPTION

FINAL STRATEGY 1 / FINAL OUTREACH 1

WHY THIS OPTION

It has the strongest evidence chain:

Autostep's core product → workflow-data clustering → need for reliable classification → Theresia's demonstrated validation/data-auditing experience → small, testable intervention.

It also avoids the biggest outreach risk: telling a technically sophisticated founder that their system has a problem when there is no public evidence that it does.

HUMAN REVIEW REQUIRED

Before sending, Theresia should verify:

- *the company need is actually worth pursuing;*
- *the contact is the right person;*
- *the email address is correct if one becomes available;*
- *the proposed intervention is realistic;*
- *the message sounds like her;*
- *she is comfortable sending it.*

Run 3 — Cignara | Input: Cignara | Website: cignara.com

Result:

FINAL OUTREACH 1

Hi Nalin,

Cignara's approach to customer support is interesting because the agent isn't just answering a question — it has to use customer context, apply company policy, take an action, and verify that action.

That makes the boundary cases interesting: incomplete customer context, conflicting rules, ambiguous eligibility, or requests where the correct behaviour is to escalate rather than act.

I'd be interested in taking one bounded workflow and deliberately stress-testing it with cases like these, then mapping where the current checks handle them well and where there may be uncovered cases.

I did something similar with an agentic BI system I built: I injected 24 issues into the workflow and caught all 24 through stress testing.

If that sounds useful, would you be open to a short conversation about one workflow that would make a good test?

FINAL STRATEGY 2

Workflow regression testing.

The focus is not on finding individual agent failures, but on creating a repeatable test layer for a workflow as its prompts, policies, retrieval or logic evolve.

FINAL OUTREACH 2

Hi Nalin,

Cignara's agents sit across customer context, knowledge, policies, workflows and actions. As those pieces evolve across enterprise deployments, one question I'd want to make easy to answer is whether a change has quietly broken an interaction that previously worked.

I'd be interested in trying a small regression-testing layer around one workflow: representative cases, deliberately difficult cases, and explicit expected outcomes that can be rerun whenever the workflow changes.

I've built reproducible ML workflows, Pytest-based systems, and an agentic system with reflection and retry logic, so the work I'd be bringing is the testing layer rather than another model.

If there's a workflow where repeatable evaluation would be useful, I'd be interested in seeing whether a small version of this adds anything.

FINAL STRATEGY 3

Confidence and grounding validation.

The focus is on testing whether the confidence/citation signals Cignara exposes correspond to actual answer quality.

FINAL OUTREACH 3

Hi Nalin,

Cignara puts unusual emphasis on grounding, confidence signals, citations and continuous evaluation. One thing I'd be curious to test is whether those signals stay aligned with correctness when a customer request is ambiguous, contradictory, or only partially supported by the available information.

A small trial could use synthetic or anonymized conversations labelled as supported, partially supported, contradictory and unanswerable, then compare the agent's confidence and citations with the actual outcome.

I've worked on model evaluation where the headline metric initially looked excellent but turned out to be misleading because of data leakage. That made me particularly interested in testing whether an evaluation signal is actually measuring what we think it is.

Would it be worth discussing one narrow evaluation set that could test this?

RECOMMENDED OPTION

FINAL OUTREACH 1 — Adversarial workflow test.

WHY THIS OPTION

It has the best combination of:

- *strongest company evidence;*
- *strongest Theresia evidence;*
- *smallest intervention;*
- *lowest assumption burden;*
- *clearest technical value;*
- *easiest demonstration of usefulness.*

Most importantly, it does not require Cignara to admit that anything is broken.

The offer is to test a core property of the product.

That distinction materially reduces the risk of sounding presumptuous.

HUMAN REVIEW REQUIRED

Before sending, Theresia should verify:

- *the company need is actually worth pursuing;*
- *the contact is the right person;*
- *the email address is correct;*
- *the proposed intervention is realistic;*
- *the message sounds like her;*
- *she is comfortable sending it.*

Because no professional email was verified for Nalin Gupta, the email version should NOT be sent until a legitimate contact address is established. A LinkedIn approach is the safer channel if no verified professional email can be found.

Run 4 — Light Anchor | Input: Light Anchor | **Website:** lightanchor.ai

Result:

FINAL OUTREACH 1

Chase,

Light Anchor is doing something more interesting than just adding agents to an e-commerce workflow: you're letting them actually make operating decisions across marketing, inventory, sourcing, support and other parts of the business, with escalation when something needs human judgment.

That makes me more interested in the failure cases than the happy path.

I'd like to take one bounded workflow and deliberately try to break it: inject bad data, ambiguous instructions, tool failures and commercially wrong-but-plausible outputs, then see which cases the current system catches.

I did a similar exercise on my own agentic BI system, where I injected 24 issues and caught all 24 through the evaluation suite. I've also uncovered data leakage in an ML evaluation by investigating results that looked too good to be true.

If you have a small anonymized or synthetic slice where this would be useful, I could turn the test into a short failure map rather than propose a larger project.

Worth a short conversation?

FINAL STRATEGY 2

Adversarial data-quality audit of one existing operational data flow.

FINAL OUTREACH 2

Chase,

Your public description of Light Anchor includes ingestion, cleansing, normalization and reconciliation, and those outputs ultimately feed agents making operational decisions.

I'd be interested in testing one small part of that data flow specifically for errors that are easy to accept because they still look valid: conflicting records, missing values, schema mismatches and semantic inconsistencies.

I've done data auditing and validation across ML pipelines at FlyRank and built end-to-end ingestion/API systems where I had to own the integration as well as the model.

If you have an anonymized slice of a messy operational dataset, I could turn the audit into a small set of discrepancy checks and concrete failure cases.

Would that be useful to investigate?

FINAL STRATEGY 3

Independent analysis of AdBench's model-ranking metrics.

FINAL OUTREACH 3

Chase,

I saw that Light Anchor is building AdBench to compare models running advertising workflows across metrics such as spend, CTR, CPC and outbound performance.

One thing I'd be curious to test is whether the "best" model changes depending on the business objective. A model that wins on clicks, for example, isn't necessarily the one you want if the downstream objective is different.

I could take a small result set and analyze where model rankings change by metric, which metrics agree, and where a benchmark result could give a misleading impression of overall performance.

I've done business-metric-driven model selection in ML projects and built an agentic BI system with reflection and automated stress testing.

If AdBench is still evolving, I'd be happy to investigate one narrow slice rather than make a larger proposal.

Open to a short call?

RECOMMENDED OPTION

FINAL STRATEGY 1 — adversarial agent evaluation.

WHY THIS OPTION

It has the best combination of:

- *strong evidence;*
- *direct relevance to Light Anchor's core thesis;*
- *demonstrated Theresia capability;*
- *small trial potential;*
- *low assumption burden;*
- *clear differentiation from generic ML outreach.*

It also does not require claiming that Light Anchor has a technical deficiency. The proposition is simply: "Let me test the system against failure cases and show you what happens."

HUMAN REVIEW REQUIRED

Before sending, Theresia should verify:

- *the company need is actually worth pursuing;*
- *Chase Kim is the right person to contact;*
- *founders@lightanchor.ai is still the correct email address;*
- *the proposed intervention is realistic for Light Anchor to permit;*
- *the message sounds like her;*
- *she is comfortable sending it.*

Run 5 — Minicor | Input: Minicor | Website: minicor.com

Result:

FINAL STRATEGY 1

Deliberately break one representative workflow in a sandbox and measure recovery behavior.

FINAL OUTREACH 1

Subject: Breaking one of Minicor's workflows

Hi Faizaan,

You describe one of Minicor's hardest problems pretty clearly: desktop automations break under real-world conditions, and the product has to recover rather than leave someone fixing every failure manually. With 2.5M executions now showing on the site, I'd be curious to see whether I could find a recovery blind spot rather than assume there is one.

If you gave me one synthetic or anonymized workflow, I'd deliberately inject a small set of realistic failures and measure which ones are detected, recovered correctly, recovered incorrectly, or missed.

I did a similar exercise on an agentic system I built: I injected 24 issues and caught all 24 through the test suite.

If that sounds useful, I'd be happy to keep it to one workflow and see whether I can find anything your current tests don't cover.

FINAL STRATEGY 2

Use existing execution traces to test whether failure diagnosis can be made more systematic.

FINAL OUTREACH 2

Subject: A question about Minicor's execution traces

Hi Faizaan,

Minicor already exposes run-level logs, step-level logs and video replay for its automations. I was curious about a narrower question: can those traces reliably distinguish "the workflow completed" from "the workflow recovered correctly"?

I'd like to take a small anonymized sample, classify the failure modes, and see where the existing evidence is enough to explain what happened versus where an engineer still has to reconstruct it manually.

My relevant work has been around making ML and agent systems testable rather than treating outputs as automatically trustworthy. On one agentic BI system, I built a fault-injection suite that caught 24/24 injected issues.

If your current traces already answer this well, that's useful to know too. If not, I could make the investigation small enough to test the idea without touching customer data.

FINAL STRATEGY 3

Create a compact synthetic regression suite around recurring automation failures.

FINAL OUTREACH 3

Subject: A small regression test for desktop automation

Hi Faizaan,

As Minicor adds workflows and legacy environments, I imagine one useful question is how to regression-test recovery without needing a full customer environment for every change.

If you're interested in exploring that, I could take one representative workflow, create controlled perturbations around it, and turn them into repeatable tests with expected outcomes. The goal wouldn't be to replace your existing testing—just to see whether a small adversarial suite can catch anything before production.

I built a similar stress-testing setup for an agentic system and caught all 24 injected issues.

If a one-workflow experiment would be useful, I could start there.

RECOMMENDED OPTION

FINAL OUTREACH 1

WHY THIS OPTION

Strategy 1 has the best evidence-to-assumption ratio.

It starts from an explicitly documented Minicor problem, uses a highly specific piece of Theresia's demonstrated experience, requires only a small experiment, and does not claim that Minicor's current engineering is deficient.

It also creates a credible way for Minicor to evaluate Theresia through actual technical work rather than through a résumé.

HUMAN REVIEW REQUIRED

Before sending, Theresia should verify:

- the company need is actually worth pursuing*
- the contact is the right person*
- the email address is correct*
- the proposed intervention is realistic*
- the message sounds like her*
- she is comfortable sending it*

In particular, she should decide whether the 2.5M execution figure is sufficiently current and whether she wants to reference it directly. The figure is currently displayed on Minicor's website.

3. Time Accounting and Time Saved

The manual version of this process would realistically take **more than four hours per company**.

This is because the research is not simply a matter of asking an AI model one question.

A manual process would involve:

- researching the company and its website;
- finding current information;
- identifying relevant people;
- checking hiring activity;
- understanding the company's technical direction;
- assessing potential role and project fit;
- synthesising the information;
- and finally turning the research into a useful written output.

Using the structured workflow, the completed research-to-drafting process took approximately **2–3 minutes per company once the workflow/prompt structure was established**.

This represents a substantial reduction in execution time.

Important qualification

I am not treating the 2–3 minute figure as the total project time.

There was significant **setup and iteration time** involved in designing the workflow, configuring the n8n nodes, testing the available models/APIs, discovering the web-search limitation, and restructuring the process into separate ChatGPT prompts.

Therefore, the honest comparison is:

Process	Approximate execution time
Manual research → drafting	4+ hours per company
Structured AI workflow	2–3 minutes per company
Initial workflow/setup	Several hours

The automation is therefore most valuable when it is **reused across multiple companies**, because the setup cost is amortised over subsequent runs.

4. Known Failure Points and Human Review

The workflow significantly reduces the mechanical research and synthesis work, but it does not remove the need for human judgement.

1. AI judgement about what is important

The model can identify information and produce a coherent synthesis, but it does not necessarily know which signals are actually the most important.

For example, it may correctly identify several hiring signals but give too much weight to a less relevant signal and not enough weight to something that a human would immediately recognise as strategically important.

2. Role-fit judgement is inherently limited

The AI can assess apparent fit based on the information provided to it, but it cannot fully determine whether a particular role, company or opportunity is actually the right fit for me.

It does not truly know what I want to work on, what type of work I want to embark on, what I would find interesting, or what constraints and preferences I would apply when making a career decision.

Therefore, its role-fit analysis is necessarily an approximation based on the available evidence.

3. Web research can still miss information

Even when web search is available, the model may fail to find an important page, recent announcement, job posting, person, or other signal.

A search result being absent does not necessarily mean that the underlying information does not exist.

4. Information can change

Company websites, job postings, team structures, products and priorities can change rapidly. A result that is accurate at the time of research may become outdated.

This means the workflow is useful precisely because it can perform fresh research, but the resulting output should still be treated as a point-in-time assessment.

5. Human verification remains necessary

Before acting on the final output, I would still manually verify the most important claims, particularly:

- current employment and job titles;
- active job postings;
- company/product claims;
- technical direction;
- important people;
- and any conclusion that would materially affect an outreach or career decision.

6. Prompt Length

Current prompts were built for n8n workflow in mind not optimized for ChatGPT chat terminal, future iterations can make the prompt less specific so that the detailed prompting does not negatively affect the output

5. Prompts / Configuration Used

Configuration = Chatgpt Project each new run was a completely new chat

Step 1 — Company Research

Purpose: Research the company and identify relevant people, hiring signals, technical direction, and evidence-backed opportunities.

You are Step 1 of a multi-stage company research and cold outreach workflow.

Your task is to conduct comprehensive, current research on the company provided below using web search and publicly available sources.

COMPANY INPUT

Company name:
{{ \$json.company_name }}

Company website:
{{ \$json.company_url }}

Optional context:
{{ \$json.optional_context }}

RESEARCH OBJECTIVE

Build a factual, evidence-based company profile that can later be used to determine:

1. What the company actually does.
2. What it is currently trying to accomplish.
3. What technical or operational problems it may currently be dealing with.
4. What it is actively hiring for.
5. Who the relevant decision-makers are.
6. Where there may be a credible opportunity for Theresia Saumu to contribute.

Do not write outreach at this stage.

RESEARCH PRIORITY

Prioritise, in order where available:

- Official company website
- Official company documentation
- Company blog/newsroom
- Official announcements
- Current job postings
- Founder and leadership profiles
- Funding announcements
- Public engineering/technical material
- Reputable reporting
- Public professional profiles

Use multiple sources when possible.

For every important claim, retain the source URL.

COMPANY INFORMATION TO RESEARCH

Research:

- Company name
- Website
- Industry
- Location
- Founded date
- Approximate age
- Approximate employee/team size
- Company stage
- Funding raised
- Most recent funding
- Investors
- Mission
- Philosophy
- Strategic direction
- Current outlook
- Products

- Services
- Customers/users
- Problems the company solves
- How its products/solutions work
- Business model if publicly available

TECHNICAL LANDSCAPE

Research where publicly identifiable:

- Technology stack
- Programming languages
- Frameworks
- Cloud/platform infrastructure
- AI/ML usage
- Data systems
- APIs
- Integrations
- Automation
- Technical architecture
- Engineering practices
- Technical challenges
- Areas of technical expansion

Do not infer a specific technology merely because it is common in the industry.

CURRENT ACTIVITY

Research recent:

- Product launches
- Product changes
- Partnerships
- Expansion
- Funding
- Leadership changes
- Strategic announcements
- Major customer activity
- Technical initiatives
- Other developments that could affect the company's current priorities

HIRING

Research current and recent hiring.

Pay particular attention to:

- Engineering roles
- AI/ML roles
- Data roles
- Backend roles
- Platform/infrastructure roles
- Product engineering roles
- Technical leadership roles
- Roles that indicate a newly emerging capability
- Repeated skills appearing across roles
- Roles that reveal current technical or operational priorities

Do not assume a hiring requirement is evidence of a problem by itself. Treat it as a signal.

POTENTIAL NEEDS

Identify potential technical or operational needs suggested by the evidence.

For every potential need, distinguish between:

FACT — directly stated or clearly documented.

STRONG INFERENCE — not directly stated but supported by multiple credible signals.

WEAK INFERENCE — plausible but insufficiently supported.

Do not present weak inferences as established company problems.

CONTACT DISCOVERY

Identify people who could plausibly be decision-makers for technical work or hiring.

Prioritise:

- CEO
- Founder
- CTO
- Technical Founder
- VP Engineering
- Head of Engineering
- Head of AI/ML
- Engineering Manager
- Technical hiring decision-maker
- Other person whose role is directly relevant to a specific opportunity identified

For every relevant person provide:

- Full name
- Current role
- Why they are relevant
- LinkedIn profile URL
- Professional email if found
- Email status:
 - VERIFIED
 - INFERRED
 - NOT FOUND
- Evidence for the contact information

EMAIL RULES

A professional email may be:

VERIFIED:

Directly found in a credible public source.

INFERRED:

Not directly published, but reasonably inferred from a documented company email pattern or other public evidence.

NOT FOUND:

No credible email could be established.

If an email is inferred, provide the likely address but explicitly label it INFERRED.

Never invent a contact from nothing.

Always provide the person's LinkedIn profile/name when a relevant profile can be identified.

OUTPUT FORMAT

Return plain text only.

Use exactly these headings:

COMPANY OVERVIEW

COMPANY PHILOSOPHY AND DIRECTION

PRODUCTS AND SOLUTIONS

TECHNICAL LANDSCAPE

CURRENT ACTIVITY AND RECENT NEWS

FUNDING AND COMPANY STAGE

TEAM AND ORGANISATION

HIRING SIGNALS

POTENTIAL TECHNICAL OR OPERATIONAL NEEDS

DECISION-MAKERS AND CONTACTS

CONTACT CONFIDENCE NOTES

SOURCE EVIDENCE

RESEARCH GAPS

Under SOURCE EVIDENCE, include the URLs supporting the important claims.

Do not perform personal fit analysis.

Do not recommend outreach.

Do not write outreach.

Step 2 — Need & Evidence Analysis

Input: Step 1 output

You are Step 2 of a multi-stage company research and cold outreach workflow.

Your task is to critically evaluate the company research and determine which potential needs are actually defensible.

Do not try to manufacture an interesting problem simply because the final workflow needs an outreach angle.

COMPANY RESEARCH

{{ \$('Step 1 Research').item.json.text }}

CLASSIFICATION STANDARD

For every important observation, distinguish between:

FACT

Directly supported by a credible source.

STRONG INFERENCE

Not directly stated but supported by multiple credible signals.

WEAK INFERENCE

Plausible but insufficiently supported.

A WEAK INFERENCE must not become the primary basis for an outreach strategy.

ANALYSE

Look especially at:

- Current products
- Current technical direction
- Current hiring
- Recent announcements
- Product complexity
- Scaling signals
- Operational changes
- New capabilities being built
- Repeated technical requirements
- Explicitly stated problems
- Signals that indicate where the company is investing

Identify 3–5 potential company needs or opportunities.

For each opportunity determine:

- What the need is
- Whether it is FACT, STRONG INFERENCE or WEAK INFERENCE
- Evidence supporting it
- Source supporting the evidence
- How current the signal appears to be
- Why the need might matter
- Whether it appears technically or operationally meaningful
- Whether a small trial could plausibly address it
- What remains uncertain

RANKING

Rank opportunities based on:

1. Evidence strength
2. Currentness
3. Technical/operational relevance
4. Potential usefulness
5. Suitability for a small trial or intervention
6. Low risk of being a false assumption

Do not optimise for what sounds impressive.

Optimise for what is defensible.

OUTPUT FORMAT

Return plain text only.

Use exactly these headings:

EVIDENCE SUMMARY

VERIFIED COMPANY FACTS

STRONGLY SUPPORTED NEEDS

WEAK OR SPECULATIVE NEEDS

TOP OPPORTUNITIES

OPPORTUNITY 1

OPPORTUNITY 2

OPPORTUNITY 3

ADDITIONAL OPPORTUNITIES

RECOMMENDED OPPORTUNITY

WHY THIS OPPORTUNITY IS DEFENSIBLE

WHAT IS STILL UNKNOWN

RESEARCH EVIDENCE

Do not write outreach.

Do not perform personal fit analysis.

Step 3 — Personal Fit Analysis

Input: Steps 1 + 2

You are Step 3 of a multi-stage company research and cold outreach workflow.

Your task is to determine where Theresia Saumu's demonstrated capabilities intersect with the company's evidence-backed needs.

Do not optimise for maximum keyword or technology overlap.

Optimise for:

- credible contribution
- practical usefulness
- specificity
- technical fit
- ability to demonstrate value through a small intervention
- evidence that Theresia has actually done something relevant

COMPANY RESEARCH

{{ \$('Step 1 Research').item.json.text }}

VALIDATED OPPORTUNITY ANALYSIS

{{ \$('Step 2 Evidence').item.json.text }}

THERESIA SAUMU — PROFESSIONAL KNOWLEDGE GRAPH

IDENTITY

Name: Theresia Saumu

Location: Kenya (Nairobi) / Remote

Current direction:

Machine Learning Engineering and LLM Systems

Target roles:

AI Engineer

Machine Learning Engineer

LLM Systems Engineer

Data Scientist

Analytics Engineer

MLOps Engineer

Work preference:

Remote and systems-oriented work involving high ownership, interesting challenges, real stakes and concrete problems to solve.

POSITIONING

Theresia is a systems-centric Machine Learning Engineer who builds useful working systems designed around what actually needs to happen to solve real-world problems.

AI is treated as a component of a useful system rather than the novelty being sold.

Engineering philosophy:

- System-first design
- Reliable, low-latency interfaces
- End-to-end ownership
- Willingness and demonstrated ability to take over integration, MLOps, APIs and testing when delivery requires it

EXPERIENCE

Machine Learning Intern — FlyRank AI

Jun 2026–Present

- Builds ML pipelines across large-scale search intelligence datasets.
- Owns pipeline lifecycle from problem framing through data auditing, baseline modelling and validation.
- Improved project accuracy from a 24% deterministic baseline to a 76% ML target.
- Applies production-oriented engineering and data governance.

Data Science Intern — AnalystLab Africa

Jul 2026–Present

- Builds production-oriented machine learning systems.
- Uses clean architecture and SOLID principles.
- Designs reproducible, robust and testable workflows.
- Focuses on maintainability and team usability.

SKILLS

Machine Learning / Modelling:

- Algorithm selection and architecture search
- Statistical preprocessing and normalisation
- Model comparison
- Business-metric-driven model selection
- Deep learning

AI / LLM:

- Prompt engineering
- Schema enforcement
- Agentic orchestration
- LangGraph
- Reflective multi-node AI systems

Data Engineering:

- API ingestion
- ETL
- SQLite
- PostgreSQL / Supabase
- SQL views
- Feature engineering

API / MLOps:

- Flask

- FastAPI
- Authentication
- CORS
- Gunicorn
- Railway
- Rate limiting
- Automated testing with Pytest

BI:

- Power BI
- DAX
- Row-Level Security
- Programmatic PBIR/TMDL compilation
- Enterprise relational modelling

Leadership:

- Led a 19-member cross-track team
- Technical coordination
- GitHub workflows
- Code reviews
- Coding standards
- Sprint management

PROJECT EVIDENCE

SolarSense

- Built a Flask API gateway serving three ML/deterministic services.
- Compressed solar viability assessment from 1–3 weeks to approximately 3 minutes.
- Covered 20 Sub-Saharan African cities.
- Took over backend/integration responsibilities when the assigned backend developer stepped back.
- Added API authentication and rate limiting.
- Deployed through Gunicorn/Railway.
- Top 15 of 70 capstone teams.

StudyGapAI

- Built solo in one week.
- Full-stack AI diagnostic platform.
- React/Tailwind frontend.
- Flask backend.
- Supabase/PostgreSQL/Auth/RLS.
- Gemini 2.0 Flash integration.
- Engineered a 287-line system prompt enforcing strict JSON output.
- Built around 3–5 second latency.
- Won 1st place at WUD AI Hackathon 2025.

Alzheimer's Classification

- Compared ResNet50, EfficientNetB0 and DenseNet121.
- Implemented robust IQR normalisation.
- Implemented block-wise fine-tuning.
- Investigated a 99.83% apparent test accuracy.
- Discovered critical data leakage.
- Rejected the misleading accuracy instead of presenting it as a valid clinical result.
- Demonstrates model evaluation, data integrity auditing and technical scepticism.
- Current rerun on validated data is still in progress; reported figures are not final.

Agentic BI Builder

- Built a LangGraph multi-node agent.
- Profiling → domain/KPI generation → SQL execution → reflection → retry → PBIR compilation.
- Reverse-engineered Power BI PBIR/TMDL structures.

- Used Pydantic to compile native Power BI project files.
- Tested across multiple datasets.
- Caught 24/24 injected issues in stress testing.
- Demonstrates agentic orchestration, deterministic/LLM hybrid architecture and self-checking systems.

Customer Churn Platform

- Built end-to-end ETL, training, API and BI platform collaboratively.
- Evaluated five ML models.
- Built FastAPI service.
- Built Pytest suites.
- Took over final model/API integration.
- Chose Random Forest based on 71% recall rather than raw accuracy.
- 7,032 records.
- Demonstrates business-driven ML, technical leadership and integration.

Financial Risk Monitor

- Built automated ingestion pipeline using yfinance and FRED.
- SQLite-based analytics.
- Forecasting with Prophet.
- VaR/CVaR.
- 10,000-path Monte Carlo simulation.
- Power BI dashboard.
- Used DAX What-If parameters to avoid slow runtime Python recalculation.
- Monitors 11 assets.
- Demonstrates data engineering, forecasting, quantitative analytics and production-oriented system design.

EDUCATION / CERTIFICATIONS

- Women Techsters Data Science & Engineering Fellowship, Tech4Dev
- Technical Lead for 19-member team
- Top 15 of 70 capstone teams
- IBM Data Science Professional Certificate
- Google Data Analytics Professional Certificate
- Two completed academic years in Electrical and Electronic Engineering & IT at University of Johannesburg
- 10 academic distinctions
- No degree awarded

IMPORTANT VERIFICATION RULES

Do not claim:

- experience that is not documented
- production experience that is not documented
- metrics not included above
- expertise solely because a technology appears in the knowledge graph
- individual ownership of a collaborative project where the contribution is not documented
- a completed degree

FIT ANALYSIS TASK

For each strong company opportunity determine:

1. What the company appears to need.
2. What evidence supports the need.
3. Which of Theresia's demonstrated capabilities are genuinely relevant.
4. Which project, role or accomplishment proves that capability.
5. What Theresia could realistically build, test, investigate, improve or validate.
6. What a small trial could look like.
7. Why that intervention would be useful.
8. What Theresia should NOT claim.
9. Whether the fit is strong enough to justify outreach.

Prefer specific interventions over generic statements such as "I can help with AI" or "I have experience in machine learning."

POSITIONING PRINCIPLE

Use this as the underlying positioning:

"Useful working systems, built around what actually needs to happen."

Do not mechanically insert this sentence into the eventual outreach.

OUTPUT FORMAT

Return plain text only.

Use exactly these headings:

FIT SUMMARY

COMPANY NEED

EVIDENCE FOR THE NEED

RELEVANT USER CAPABILITIES

PROJECT / EXPERIENCE EVIDENCE

POSSIBLE CONTRIBUTION

POSSIBLE SMALL TRIAL

WHY THE CONTRIBUTION MATTERS

WHAT THE USER SHOULD NOT CLAIM

FIT STRENGTH

BEST OPPORTUNITY

ALTERNATIVE OPPORTUNITIES

KEY UNCERTAINTIES

Do not write outreach.

Step 4 — Outreach Strategy

Input: Steps 1–3

You are Step 4 of a multi-stage company research and cold outreach workflow.

Your task is to design exactly THREE genuinely different outreach strategies for approaching the company.

Do not write the actual outreach messages yet.

COMPANY RESEARCH

{{ \$('Step 1 Research').item.json.text }}

EVIDENCE / OPPORTUNITY ANALYSIS

{{ \$('Step 2 Evidence').item.json.text }}

PERSONAL FIT ANALYSIS

{{ \$('Step 3 Fit').item.json.text }}

THERESIA SAUMU — POSITIONING

Theresia is a systems-centric Machine Learning Engineer who builds useful working systems designed around what actually needs to happen to solve real-world problems.

Core positioning:

"Useful working systems, built around what actually needs to happen."

The emphasis is on the useful system and the work that needs to happen, not AI as a standalone novelty.

VOICE

Curious, analytical, direct, conversational, precise, grounded.

Voice principles:

- Don't hedge unnecessarily.
- State the conclusion and the reasoning.
- Don't use filler or generic enthusiasm.
- Make caveats specific.
- Prefer concrete observations over vague claims.
- If a claim is strong, support it with the actual reasoning.
- Sound like a person thinking through a problem, not a person performing expertise.

OUTREACH RULES

- Lead with the company's situation rather than Theresia's biography.
- Demonstrate research through a concrete observation.
- Focus on a real company need.
- Use evidence rather than invented pain points.
- Do not turn the message into a CV.
- Do not ask directly for a job.
- Do not use generic praise.
- Do not sound like a sales pitch.
- Do not exaggerate.
- Use Theresia's experience as evidence, not as the focus.
- Prefer a concrete trial, investigation, test, validation or build.
- Email is preferred when a legitimate professional email exists.
- Inferred email addresses must be clearly labelled INFERRED.
- Never fabricate contact information.
- Always provide LinkedIn profile/name where available.

REFERENCE OUTREACH EXAMPLES

The examples below represent Theresia's current experimentation with outreach positioning and tone.

They are NOT gold-standard templates.

Do not copy their wording, structure, sentence patterns or rhetorical devices.

Use them only to understand tendencies in her current approach.

1. Hi Marshall, Your Founding Engineer post mentions an onboarding agent that ingests a customer's historical data and task-mining video and turns it into their custom agent and ERP. The part I'd worry about isn't the model, it's the ingestion: RFQs, job costs, and drawings arriving in enough different formats that "train on this" quietly means "clean this first," and a bad clean produces a confidently wrong quote instead of an obviously broken one. Send me one messy, anonymized customer dataset, RFQ emails, a job-cost sheet, whatever's actually painful, and I'll turn it into a validated, model-ready pipeline with the discrepancy checks and margin-floor logic already built in, the same shape

of problem I solved shipping SolarSense from an unscoped problem to a live deployment across 20 cities, and absorbed the backend for when the assigned engineer stepped back. You'd be judging a working pipeline against your own data, not a pitch. Worth a short call to pick the dataset?

2. Hi Nathan, Magic Monitor scores every extracted data point for confidence. My guess is the hard part isn't generating that score, it's knowing whether the score itself is right, especially now that Harbor is running full trials and a wrong "trust this" costs more than it used to. I'd like to try to break it. Give me a slice of anonymized or synthetic source documents, and I'll deliberately inject the kinds of errors extraction actually misses (ambiguous handwriting, transposed values, a field that's technically present but means something else in context) and report back where the confidence score agrees with reality and where it doesn't. I ran this exact exercise on my own system before trusting it: built a fault injection suite, planted 24 realistic errors into already-validated output, and used the results to find three bugs no amount of reading the code would have surfaced. I'd rather show you the gap than describe my resume. Open to a short call to figure out if this is worth a real trial?

3. (not relevant experience but is someone else's correct outreach strategy that got positive response) Hi Elana,

So, I have a bit of a different request.

My Graduate Programme lets me take up to ~8 months off to go work in industry. I know your organisation does not have any ESG or sustainability position open, which is perfect because I'd like to come help and be more involved than any intern would.

Is it possible to create a "ESG and sustainability analyst" role where I'd get to help you with developing your decarbonization action plans, setting net-zero targets, and writing ESG disclosures and reports?

I'm happy to work for free if it means I get to learn.

If there's anything from helping you engage with the financial sector and industry to align your operations with regulatory and stakeholders needs, enhancing your resilience against physical and transitional risks, or just chewing glass in some way to prove myself, I'd love to do it.

Thanks,
Abiola
Your ESG and sustainability analyst

STRATEGY TASK

Generate exactly THREE distinct strategies.

The strategies must differ in the proposed intervention, not merely in wording.

Possible strategy types include:

- Testing or deliberately trying to break something
- Building a small technical slice
- Diagnosing a technical or operational problem
- Validating an assumption
- Improving an existing workflow
- Prototyping a capability
- Investigating a bottleneck

Only use a strategy if it makes sense for the specific company and evidence.

For each strategy provide:

- Strategy name
- Core idea
- Company need addressed
- Evidence
- Proposed intervention
- What Theresa would actually do
- Relevant credibility/evidence from her background
- Ideal recipient
- Why that recipient would care
- Likely objection
- Biggest assumption
- Why this strategy is meaningfully different from the other two

Do not draft messages.

OUTPUT FORMAT

Return plain text only.

Use exactly these headings:

STRATEGY OVERVIEW

STRATEGY 1

CORE IDEA

COMPANY NEED

EVIDENCE

PROPOSED INTERVENTION

USER CREDIBILITY

IDEAL RECIPIENT

WHY THEY MIGHT CARE

LIKELY OBJECTION

KEY ASSUMPTION

STRATEGY 2

CORE IDEA

COMPANY NEED

EVIDENCE

PROPOSED INTERVENTION

USER CREDIBILITY

IDEAL RECIPIENT

WHY THEY MIGHT CARE

LIKELY OBJECTION

KEY ASSUMPTION

STRATEGY 3

CORE IDEA

COMPANY NEED

EVIDENCE

PROPOSED INTERVENTION

USER CREDIBILITY

IDEAL RECIPIENT

WHY THEY MIGHT CARE

LIKELY OBJECTION

KEY ASSUMPTION

STRATEGY DIFFERENTIATION

RECOMMENDED STRATEGY

Step 5 — Draft Outreach

Input: Steps 1, 3 and 4

You are Step 5 of a multi-stage company research and cold outreach workflow.

Your task is to turn the three approved strategic approaches into three distinct outreach drafts.

Do not invent facts.

COMPANY RESEARCH

{{ \$('Step 1 Research').item.json.text }}

PERSONAL FIT ANALYSIS

{{ \$('Step 3 Fit').item.json.text }}

OUTREACH STRATEGIES

{{ \$('Step 4 Strategy').item.json.text }}

CONTACTS

Use the decision-makers and contact information identified in Step 1:

{{ \$('Step 1 Research').item.json.text }}

THERESIA SAUMU — POSITIONING

Theresa is a systems-centric Machine Learning Engineer who builds useful working systems designed around what actually needs to happen to solve real-world problems.

Core positioning:

"Useful working systems, built around what actually needs to happen."

This is a positioning principle, not a phrase that should automatically appear in the message.

VOICE

Curious, analytical, direct, conversational, precise, grounded.

- Don't hedge unnecessarily.
- State the conclusion and reasoning.
- Don't use filler or generic enthusiasm.
- Make caveats specific.
- Prefer concrete observations.
- Support strong claims with reasoning.
- Sound like a person thinking through a problem, not performing expertise.

OUTREACH RULES

- Lead with the company's situation.
- Demonstrate actual research through specifics.
- Focus primarily on their need.
- Identify a concrete problem or opportunity.
- Propose a concrete intervention.
- Use Theresia's background only as supporting evidence.
- Do not write a CV in paragraph form.
- Do not ask directly for a job.
- Do not use generic praise.
- Do not sound like a sales pitch.
- Do not exaggerate.
- End with a low-friction next step.
- Prefer email when a professional email exists.
- If no suitable professional email exists, prepare LinkedIn outreach.
- An inferred email must be labelled INFERRED outside the message.
- Never fabricate a contact.

REFERENCE EXAMPLES

These are examples of current experimentation only, not templates:

[PASTE THE TWO CURRENT OUTREACH EXAMPLES HERE]

WRITING TASK

Write exactly three messages, one based on each strategy.

The three messages must remain meaningfully different in approach.

Each message should:

- open with something specific to the company or recipient
- show understanding rather than saying "I researched you"
- identify an evidenced need
- propose something specific Theresia could do
- make the proposed intervention feel practical
- use one or two pieces of relevant evidence from Theresia's experience where useful
- avoid unnecessary credentials
- avoid generic self-introductions
- avoid excessive flattery
- avoid corporate language
- avoid exaggerated certainty
- avoid sounding like a recruitment pitch
- avoid sounding transactional
- end with a simple next step

Do not introduce facts that do not exist in the supplied information.

CONTACT SELECTION

Select the strongest relevant decision-maker based on the proposed opportunity.

Provide their:

- name
- role
- LinkedIn
- email
- email status
- reason they are the appropriate recipient

OUTPUT FORMAT

Return plain text only.

Use exactly these headings:

CONTACT RECOMMENDATION

CONTACT NAME

CONTACT ROLE

LINKEDIN

EMAIL

EMAIL STATUS

WHY THIS CONTACT

STRATEGY 1 DRAFT

STRATEGY 2 DRAFT

STRATEGY 3 DRAFT

DRAFT NOTES

UNRESOLVED QUESTIONS

Step 6 — Adversarial Review

Input: Steps 1–5

You are Step 6 of a multi-stage company research and cold outreach workflow.

You are the adversarial reviewer.

You did NOT write the outreach drafts.

Your job is to find reasons the drafts should NOT be sent and identify exactly what must change before they are presented to Theresa.

COMPANY RESEARCH

{{ \$('Step 1 Research').item.json.text }}

EVIDENCE / OPPORTUNITY ANALYSIS

{{ \$('Step 2 Evidence').item.json.text }}

FIT ANALYSIS

{{ \$('Step 3 Fit').item.json.text }}

OUTREACH STRATEGIES

{{ \$('Step 4 Strategy').item.json.text }}

OUTREACH DRAFTS

{{ \$('Step 5 Draft').item.json.text }}

THERESIA SAUMU — KNOWLEDGE GRAPH

Theresa is a systems-centric Machine Learning Engineer who builds useful working systems designed around what actually needs to happen to solve real-world problems.

Relevant evidence includes:

- FlyRank AI ML internship: owns ML pipeline lifecycle from problem framing through auditing, modelling and validation; improved project accuracy from 24% deterministic baseline to 76% ML target.
- AnalystLab Africa data science internship: production-oriented ML, clean architecture, SOLID principles, reproducible and testable workflows.
- SolarSense: full Flask gateway integrating ML/deterministic services; compressed solar viability assessment to approximately 3 minutes; took over backend/integration responsibilities; Railway/Gunicorn deployment; 20 cities.
- StudyGapAI: built solo in one week; full-stack AI diagnostic system; React/Tailwind, Flask, Supabase, Gemini; strict structured-output prompt engineering; 1st place WUD AI Hackathon 2025.
- Alzheimer's project: compared deep learning architectures, corrected scanner artefacts, investigated suspicious performance and uncovered critical data leakage rather than presenting misleading results as valid.
- Agentic BI Builder: LangGraph multi-node agent with profiling, LLM reasoning, SQL execution, reflection, retry and Power BI compilation; caught 24/24 injected issues.
- Customer Churn: end-to-end ML/API/BI platform; model comparison; Pytest; FastAPI; technical leadership; selected model based on 71% recall.
- Financial Risk Monitor: yfinance/FRED ingestion, SQLite, Prophet, risk analytics, Monte Carlo and Power BI.

Do not assume anything beyond the documented evidence.

POSITIONING

"Useful working systems, built around what actually needs to happen."

VOICE

Curious, analytical, direct, conversational, precise, grounded.

OUTREACH PRINCIPLES

- Lead with their situation.
- Be specific.
- Show research through observations.
- Avoid generic praise.
- Avoid job-application framing.
- Avoid sales language.
- Avoid unsupported assumptions.
- Use experience as evidence rather than the subject.
- Prefer useful concrete interventions.
- Keep the ask low friction.

REVIEW CRITERIA

Evaluate every draft for:

1. COMPANY UNDERSTANDING

Does the message demonstrate real understanding?

2. EVIDENCE

Is the claimed company need actually supported?

3. PAIN VALIDITY

Is it likely to be an actual concern or is the workflow manufacturing a pain point?

4. FIT

Does Theresa have credible evidence to address it?

5. SPECIFICITY

Could this message be sent to 50 unrelated companies?

6. POSITIONING

Does it position Theresa around useful work rather than asking for employment?

7. CREDIBILITY

Are all claims about Theresa accurate and supported?

8. VOICE

Does it sound curious, analytical, direct, conversational, precise and grounded?

9. LENGTH

Does every sentence earn its place?

10. ASK

Is the proposed next step reasonable?

11. RECIPIENT

Is this actually the right person?

12. RISK

Does anything sound presumptuous, creepy, overly familiar, salesy, inflated or fabricated?

13. STRATEGIC DIFFERENTIATION

Do the three drafts actually represent different approaches?

For each issue, explain:

- What is wrong
- Why it matters
- What specifically should change

Do not rewrite the messages.

OUTPUT FORMAT

Return plain text only.

Use exactly these headings:

OVERALL ASSESSMENT

STRATEGY 1 REVIEW

STRATEGY 1 STRENGTHS

STRATEGY 1 PROBLEMS

STRATEGY 1 REQUIRED CHANGES

STRATEGY 2 REVIEW

STRATEGY 2 STRENGTHS

STRATEGY 2 PROBLEMS

STRATEGY 2 REQUIRED CHANGES

STRATEGY 3 REVIEW

STRATEGY 3 STRENGTHS

STRATEGY 3 PROBLEMS

STRATEGY 3 REQUIRED CHANGES

CONTACT REVIEW

VOICE REVIEW

POSITIONING REVIEW

EVIDENCE REVIEW

HIGHEST PRIORITY CHANGES

SEND / DO NOT SEND RECOMMENDATION

Step 7 — Revision & Final Package

You are Step 7, the final revision and packaging stage of a multi-stage company research and cold outreach workflow.

Your task is to revise the three outreach drafts using the adversarial review and return a final package for Theresia to evaluate.

Do not send anything to the prospect.

COMPANY RESEARCH

{{ \$('Step 1 Research').item.json.text }}

EVIDENCE / OPPORTUNITY ANALYSIS

{{ \$('Step 2 Evidence').item.json.text }}

FIT ANALYSIS

{{ \$('Step 3 Fit').item.json.text }}

OUTREACH STRATEGIES

{{ \$('Step 4 Strategy').item.json.text }}

ORIGINAL DRAFTS

{{ \$('Step 5 Draft').item.json.text }}

ADVERSARIAL REVIEW

{{ \$('Step 6 Review').item.json.text }}

THERESIA SAUMU — POSITIONING

Theresia is a systems-centric Machine Learning Engineer who builds useful working systems designed around what actually needs to happen to solve real-world problems.

Core positioning:

"Useful working systems, built around what actually needs to happen."

Treat this as the underlying positioning principle rather than forcing the sentence into the outreach.

VOICE

Curious, analytical, direct, conversational, precise, grounded.

- Don't hedge unnecessarily.
- State the conclusion and the reasoning.
- Don't use filler or generic enthusiasm.
- Make caveats specific.
- Prefer concrete observations over vague claims.
- Support strong claims with actual reasoning.
- Sound like a person thinking through a problem, not performing expertise.

OUTREACH RULES

- Lead with the company's situation.
- Demonstrate research with specifics.
- Focus on their need.
- Use evidence rather than invented problems.
- Use Theresia's experience as proof, not as the message's subject.
- Do not turn the message into a CV.
- Do not ask directly for a job.
- Avoid generic company praise.
- Avoid corporate buzzwords.
- Avoid sales language.
- Avoid exaggerated claims.
- Propose something concrete and useful.
- Prefer a short trial, investigation, test, validation or build.
- Keep the next step low friction.
- Email is preferred.
- Inferred email addresses must be labelled INFERRED.
- Never fabricate contact details.

REVISION RULES

Apply valid reviewer feedback.

Do not blindly follow reviewer recommendations if they conflict with the evidence.

Remove unsupported claims.

Do not introduce new company facts.

Do not introduce new claims about Theresia.

Keep the three strategic approaches genuinely distinct.

Tighten the writing.

Every sentence must earn its place.

The result should feel like a technically capable person noticed something specific and has a concrete way to be useful, not someone trying to convince the company to hire them.

FINAL OUTPUT FORMAT

Return plain text only.

Use exactly these headings:

COMPANY

RECOMMENDED CONTACT

CONTACT ROLE

LINKEDIN PROFILE

EMAIL

EMAIL STATUS

WHY THIS CONTACT

CORE OPPORTUNITY

EVIDENCE FOR THE OPPORTUNITY

WHAT REMAINS UNCERTAIN

FINAL STRATEGY 1

FINAL OUTREACH 1

FINAL STRATEGY 2

FINAL OUTREACH 2

FINAL STRATEGY 3

FINAL OUTREACH 3

RECOMMENDED OPTION

WHY THIS OPTION

HUMAN REVIEW REQUIRED

SOURCE LINKS

HUMAN REVIEW REQUIRED must explicitly remind Theresia to verify:

- the company need is actually worth pursuing
- the contact is the right person
- the email address is correct
- the proposed intervention is realistic
- the message sounds like her
- she is comfortable sending it

Do not send the prospect outreach.