

Survive the Crit

1. Peer Review Feedback

Reviewer: Latifah Bashir

Review date: 30 August 2026

I submitted the live portfolio to Latifah and gave her my proof statement so she could judge the site against what it was supposed to prove.

I asked her the two required questions first:

1. **In ten seconds, what do I do?**
2. **Would you believe I'm good at it?**

What she said

The proof statement works. Latifah said she liked the proof statement and felt that it served its purpose. She also specifically liked the favicon and the bottom CTA.

The ten-second test did not pass. Her main criticism was that the domain does not land quickly enough. The headline is:

“Useful working systems, built around what actually needs to happen.”

Her point was that this could describe several different types of people. It could be a product manager, an operations person, or a generalist engineer. She only learns that I specifically work in Applied AI from the third paragraph.

Her conclusion was that **Applied AI / ML needs to appear in the headline itself**, so the first line can communicate what I do without requiring the reader to continue through the page.

The work itself was credible. When asked whether she would believe I was good at what I do, Latifah said yes. More importantly, when she considered the question as if she did not already know me, she still pointed to the project previews as evidence. The Agentic BI Builder has an in-motion demo, and the Financial Risk Monitor has a concrete scenario dashboard. Those made the work look real and credible rather than just being claims.

The deeper feedback

Latifah also questioned whether my proof statement was concrete enough.

Her voice-note feedback made me realise that the current statement is stronger as a description of **my method** than as a description of **the thing I can actually do**.

I say that I start with the problem and make technical decisions around what needs to happen. That explains how I work, but it does not state the capability as directly as it could.

Her suggestion was to make the statement more concrete, along the lines of taking **messy problems and turning them into AI systems**.

I think this is a useful distinction. The portfolio is trying to prove that I can take a real problem through to a working system. The proof statement should therefore make that capability easier to see, rather than only explaining my approach to getting there.

2. Must-Fix vs Nice-to-Have

I reviewed the feedback against the actual job of the portfolio rather than treating every suggestion as equally important. I additionally did an AI review of the portfolio. Thus, the below items are as a result of the peer and AI review.

The portfolio has one job: **show an early-stage founder or technical leader that I can take a real problem, make the necessary technical decisions, and turn it into a useful working system**. The main action is to get someone to give me a scoped problem for a short trial project or build sprint.

Must-fix

1. Make the domain clear in the headline

Why it is a must-fix:

The ten-second test currently fails. The reader can see that I build “useful working systems,” but cannot immediately tell that the work is specifically Applied AI / ML.

This directly affects the portfolio's positioning, so it needs to be fixed.

Action:

Bring **Applied AI / ML** into the headline or immediately adjacent to it.

2. Make the capability in the proof statement more concrete

Why it is a must-fix:

The current proof statement explains my approach well, but the critique exposed a gap between describing **how I think** and clearly stating **what I can do**.

The portfolio needs to prove that I can turn a real problem into a working system. If the central claim is too abstract, the projects have to do too much work to explain it.

Action:

Refine the proof statement so the capability is concrete without turning it into generic AI-engineer marketing language.

3. Make the outcome of the project previews easier to see

Why it is a must-fix:

The AI review identified a second clarity issue: the project cards explain the problem but do not make the result immediately visible.

This matters because the homepage is supposed to provide enough evidence to make the claim credible before asking the visitor to go deeper. The cases themselves already follow the structure of **the problem → what I did and decided → what came of it**, so the homepage teasers should give the reader a clearer signal of what actually came out of each project.

Action:

Add a concise result/outcome to the project previews rather than restructuring the case studies.


Nice-to-have / Later

1. Move the full CTA higher on the page

The AI review suggested moving the full “Give me one problem you'd like solved by next week” block closer to the hero, or adding a shorter version beneath the headline since it is a strong part of the portfolio.

3. Evidence of must-fixes addressed

Before Fixes Screenshots



01

Agentic BI Builder

Starting a BI project means profiling the data, deciding what matters to the business, writing the queries, and building the report. That setup gets redone from scratch for every new business.

[Read the case study →](#)



02

Financial Risk Monitor

Investors don't have one place to see portfolio risk alongside the macro conditions around it. That makes questions like "what happens if rates move?" slow to answer.

[Read the case study →](#)

03

SolarSense

A proper solar feasibility study takes 1 to 3 weeks of consultant time. For SMEs across Sub-Saharan Africa, that cost can put the decision out of reach before they even know whether solar makes sense.

[Read the case study →](#)



Your Diagnostic Report

Comprehensive Analysis

Diagnostic Summary

Your performance indicates significant knowledge gaps across multiple topics, particularly in engineering, calculus, and geometry. You need to focus on building a strong foundation in these areas before tackling more complex problems. Consistent study and seeking help when needed will be crucial for your improvement.

Overall Performance



AIIMS Score Projection



04

StudyGapAI

Students can study hard and still fail because they're reinforcing what they already know while the actual gap causing the failure goes undiagnosed. A useful tool can't just grade the answer, it has to work out why the wrong answer happened.

[Read the case study →](#)

05

Alzheimer's Disease Classification

I built an Alzheimer's image classifier intended to assist with timely flagging, but a working classifier isn't enough. I needed to know which model to trust and whether its performance was actually real.

[Read the case study →](#)

DenseNet121 Test Set Confusion Matrix



After Fixes Screenshots

THERESIA

Applied AI/ML Engineer

ProjectsResumeContact

I build **systems** that connect data, models, and interfaces into **useful solutions.**

Applied AI & Machine Learning Engineer

ML Engineer Intern @ FlyRank AI

Data Science Intern @ AnalystLab Africa

I start with the problem, not the technology. I want to understand what needs to happen, what would make the solution useful, and what needs to be true for it to work. Then I choose the technical approach and build from there. The clearest example is my Agentic BI Builder, where I built the system end to end around the actual work involved in starting a BI project.

I'm now looking to apply that skill to real problems at early-stage companies. I want to work on problems that are still being figured out, build the system around what actually needs to happen, and iterate until it is actually useful.

[View my work →](#)



THERESIA SAUMU

Your Diagnostic Report

Generated on 11/10/2023

Diagnostic Summary

Your performance indicates significant knowledge gaps across multiple topics, particularly in trigonometry, calculus, and geometry. You need to focus on building a strong foundation in basic mathematical concepts before tackling more complex problems. Consistent study and seeking help when needed will be critical for your improvement.

Overall Performance

7% accuracy

Actual Questions	15
Correct answers	1
Wrong Questions	14/15
Used per Question	5s

JAMB Score Projection

26

Scored your diagnostic performance

StudyGapAI

Students can study hard and still fail because they're reinforcing what they already know while the actual gap causing the failure goes undiagnosed. A useful tool can't just grade the answer. It has to work out why the wrong answer happened.

I built a live diagnostic tool that traces wrong answers back to prerequisite gaps, then turns those gaps into a 6-week study plan. It is deployed and currently in use.

[Read the case study →](#)

Alzheimer's Disease Classification

I built an Alzheimer's image classifier intended to assist with timely flagging, but a working classifier isn't enough. I needed to know which model to trust and whether its performance was actually real.

I compared three model architectures and audited the dataset as part of model selection. That audit exposed augmented duplicate scans across the train and test split, so I rejected the original result rather than reporting a leaked performance number.

[Read the case study →](#)

DenseNet121 Test Set Confusion Matrix

misdiagnosed	1046	0	0	1
misdiagnosed	0	100	0	0
misdiagnosed	0	0	104	0
no misdiagnosis	2	0	2	976

99.83% test accuracy

Important note
Repeated augmented duplicate scans are present across the train and test split. Results are biased and not reported as a valid outcome.


01

Agentic BI Builder

Starting a BI project means profiling the data, deciding what matters to the business, writing the queries, and building the report. That setup gets redone from scratch for every new business.

I built a working pipeline that turns raw business data into a Power BI project, with profiling, KPI inference, validation, and report generation handled end to end.

[Read the case study →](#)



02

Financial Risk Monitor

Investors don't have one place to see portfolio risk alongside the macro conditions around it. That makes questions like "what happens if rates move?" slow to answer.

I built a portfolio risk tool that combines market and macro data, risk calculations, forecasts, and interactive scenario analysis in a 4-page Power BI dashboard.

[Read the case study →](#)

03

SolarSense

A proper solar feasibility study takes 1 to 3 weeks of consultant time. For SMEs across Sub-Saharan Africa, that cost can put the decision out of reach before they even know whether solar makes sense.

I built the assessment engine behind the tool, connecting energy consumption and generation models with a deterministic ROI engine through a shared API gateway. The result cut assessment time from 1 to 3 weeks to about 3 minutes.

[Read the case study →](#)

